

Reproduction of Paper: Spike-driven Transformer

Project of AI3610 Brain-inspired Intelligence, 2023 Fall, SJTU

Xiangyuan Xue, Shengmin Yang, Yi Ai

December 26, 2023

Group 5

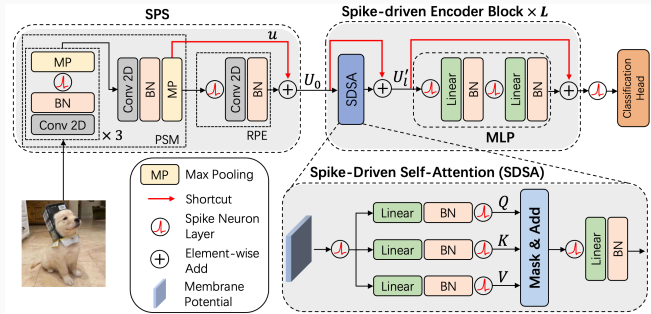
Introduction

Paper Highlight

- Spikformer proposes a spiking version of transformer, but it fails to exploit the full energy-efficiency potential of the spike-driven paradigm combined with self-attention.
- Spike-driven Transformer proposes a novel attention mechanism which only consists of mask and sparse addition operations with little energy consumption.
- Spike-driven Transformer is claimed to achieve the state-of-the-art performance on multiple common datasets and be friendly to neuromorphic hardware.

Method

Network Architecture: Overview



- Spike-driven Transformer includes 3 parts: spiking patch splitting, spike-driven encoder and classification head.

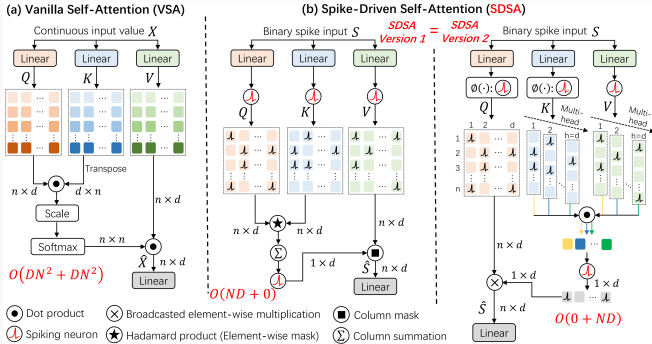
Network Architecture: Spiking Patch Splitting (SPS)

- The spiking patch splitting converts a 2D image sequence $I \in \mathbb{R}^{T \times C \times H \times W}$ into a membrane potential embedding.

$$\begin{aligned}u &= \text{PSM}(I), & u &\in \mathbb{R}^{T \times N \times D} \\s &= \mathcal{SN}(u), & s &\in \mathbb{R}^{T \times N \times D} \\ \text{RPE} &= \text{BN}(\text{Conv2d}(s)), & \text{RPE} &\in \mathbb{R}^{T \times N \times D} \\U_0 &= u + \text{RPE}, & U_0 &\in \mathbb{R}^{T \times N \times D}\end{aligned}$$

- PSM represents the patch splitting module. RPE represents the relative position embedding. $\mathcal{SN}(\cdot)$ denotes the spike neuron layer.

Network Architecture: Spike-Driven Self-Attention (SDSA)



- SDSA is different from the vanilla self-attention (VSA).

Network Architecture: Spike-Driven Self-Attention (SDSA)

- VSA is formulated as:

$$\text{VSA}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d}} \right) V$$

- SDSA is formulated as:

$$\text{SDSA}(Q, K, V) = Q_S \otimes \mathcal{SN}(\text{SUM}_c(K_S \otimes V_S))$$

where \otimes represents the Hadamard product. $\text{SUM}_c(\cdot)$ represents the column-wise summation.

Network Architecture: Spike-driven Encoder Block (SEB)

- The spike-driven encoder block follows the classical structure.

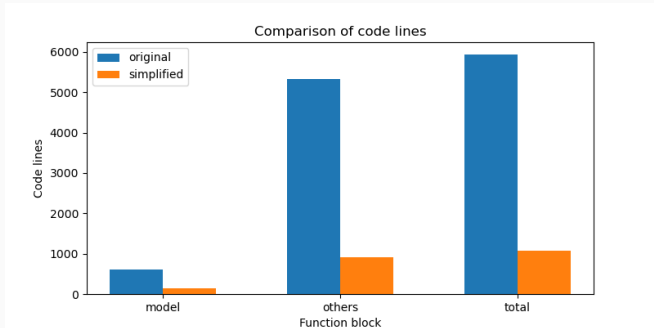
$$\begin{aligned} S_0 &= \mathcal{SN}(U_0), & S_0 &\in \mathbb{R}^{T \times N \times D} \\ U'_l &= \text{SDSA}(S_{l-1}) + U_{l-1}, & U'_l &\in \mathbb{R}^{T \times N \times D} \\ S'_l &= \mathcal{SN}(U'_l), & S'_l &\in \mathbb{R}^{T \times N \times D} \\ S_l &= \mathcal{SN}(\text{MLP}(S'_l) + U'_l), & S_l &\in \mathbb{R}^{T \times N \times D} \end{aligned}$$

- The classification head utilizes a global average-pooling and a fully connected layer to output the prediction.

Method Innovation: Simplified Implementation

- The original implementation contains lots of redundant code and applies an outdated framework, lacking modern features.
- We use the latest version of **spikingjelly**, which supports multi-step propagation, thus significantly improving the computing efficiency.
- We utilize **einops** for powerful and flexible tensor operations to simplify the code and enhance the readability.

Method Innovation: Simplified Implementation



- We implement all the function blocks within $1/6$ code lines.

Method Innovation: Transfer Learning

- Training SNNs on a GPU-based platform is extremely expensive compared to ANNs.
- We pretrain a model on the large-scale dataset to capture generic features and patterns, and then finetune it on the smaller datasets, which allows the model to adapt quickly.
- Transfer Learning can slightly improve the performance and significantly raise the training efficiency.

Experiments

Experiment Setup: Dataset

- We evaluate our model on both static datasets including CIFAR-10 and CIFAR-100 and neuromorphic datasets including CIFAR-10-DVS and DVS-128-Gesture.
- Considering the limited time and computing resources, we use the ImageNet-1k dataset only for a few rounds of pretraining instead of complete training.

Experiment Setup: Implementation Details

- We mainly use Spike-driven Transformer with 4 encoder blocks, 8 heads and 512 channels. The model is trained for 300 epochs on a single NVIDIA Tesla V100 GPU.
- We use the AdamW optimizer with a learning rate of 0.0005. The learning rate is warmed up for 10 epochs and then decayed by cosine annealing for 40 epochs.
- The batch size is 32 for static datasets and 16 for neuromorphic datasets. The number of timesteps is 4 for static datasets and 16 for neuromorphic datasets.


Result Analysis

- With only simple data augmentation employed, the metrics on all the datasets are close to those in the paper.

| Dataset | #Class | #Timestep | Accuracy |
|-----------------|--------|-----------|----------|
| CIFAR-10 | 10 | 4 | 92.90% |
| CIFAR-100 | 100 | 4 | 72.44% |
| CIFAR-10-DVS | 10 | 16 | 75.00% |
| DVS-128-Gesture | 11 | 16 | 98.96% |

Experiments


Result Analysis



- The training curves on the CIFAR-10, CIFAR-100, CIFAR-10-DVS and DVS-128-Gesture datasets show stable trends influenced by the adjustment of the learning rate.

Experiments

Transfer Learning



- With transfer learning, the model can achieve satisfactory performance within several epochs of finetuning.

| Dataset | Standard | Finetune |
|-----------|----------|---------------|
| CIFAR-10 | 89.56% | 91.25% |
| CIFAR-100 | 67.98% | 70.90% |

Conclusion

Conclusion

- We successfully reproduce the paper and verify the correctness of the paper.
- We utilize the latest frameworks to simplify the implementation and introduce the method of transfer learning to improve training efficiency.
- We conduct experiments to prove the effectiveness of our method innovation.

Thank You!